Tensor Lab: Training the Next Generation of Clinical AI Leadership

An Information Session for Medical Students

Prepare for residency. Accelerate your research.

← **Attendance!**

Our Mission

Train future physicians and AI leaders to conduct reproducible, high-impact clinical research that translates from bench to bedside.

Our Vision

Build a future where AI-integrated clinical practice becomes the standard, and physician-leaders work alongside AI experts to drive innovation.

Why Tensor Lab?

Why AI Integration Matters: The Urgent Need for Tensor Lab

The AI Revolution

AI is fundamentally reshaping clinical practice.

Most physicians lack the technical expertise to understand, validate, and deploy AI responsibly.

The Critical Gap

Severe shortage of physicians who bridge AI and clinical medicine.

AI researchers lack clinical context; clinicians lack AI expertise.

Why Tensor Lab Exists

→ **Cultivate physicians and scientists** who speak both languages

→ **Produce rigorous research** that translates AI from lab to bedside

→ **Lead the transition** into the new era of AI-integrated medicine

Meet the Directors!

UCSF School of Medicine

UNIVERSITY of MARYLAND SCHOOL OF MEDICINE

Nathan Robbins, chapter CEO

Enrique Vasquez, chapter COO

Elisa Danthinne, chapter CTO

Matt Allen

CO-FOUNDER & EXECUTIVE DIRECTOR

University of California, San Francisco School of Medicine

Aaron Ge

CO-FOUNDER & TECHNICAL DIRECTOR

University of Maryland School of Medicine

Aarushi Negi, chapter co-CEO

Safi Rifai, chapter co-CEO

Skylar Chan, chapter CTO

Gavin Shu

STRATEGIC DIRECTOR

University of California, San Francisco School of Medicine

Angie Lee

BUSINESS DIRECTOR

University of Maryland School of Medicine

Chy Murali

OPERATIONAL DIRECTOR

University of Maryland School of Medicine

Dariush Aligholizadeh, chapter co-COO

Eric Kim, chapter co-COO

The Clinical Gap: Where Medical Students Make the Difference

The Problem

Many AI projects lack clinical relevance and real-world application.

Brilliant algorithms don't translate to better patient care without clinical context.

The Solution

Medical students provide the clinical input and context.

You transform AI research into clinically significant, impactful work.

What Medical Students Bring

→ **Clinical Perspective:** Frame AI problems in ways that matter to patient care

→ **Domain Expertise:** Understand clinical workflows and real-world constraints

Enhance Your Required Summer Research: AI-Focused Support

The Reality

As a first-year medical student, you're probably doing summer research anyway.

Tensor Lab is an enhancement to that experience—not an additional burden.

Faculty Connections

- Access to top-tier mentors
- Institutional leaders in your field
- Networking opportunities
- Career guidance

Coding Support

- Learn AI & technical skills
- Hands-on programming mentorship
- Code review & feedback
- Technical infrastructure

AI Slant

- Transform your project with AI
- Clinical + technical expertise
- Publication-ready research
- Competitive advantage

The Bottom Line

Tensor Lab is an additional resource for your summer research if you want to give your project an AI focus and clinical significance.

How does it work?

The Triad Model: Core Research Unit

Each research project is led by three complementary roles working together toward publication-ready, clinically significant research.

ML Fellow

Technical Lead

- Lead technical design & coding
- Ensure reproducibility
- Own day-to-day execution
- Present at conferences

Med Mentor

Clinical Guide

- Provide clinical framing
- Guide project scoping
- Ensure clinical rigor
- Support manuscript prep

Faculty Mentor

Principal Investigator

- Provide clinical question
- Approve methodology
- Ensure IRB compliance
- Support publication

The Triad Model: Core Research Unit

Each research project is led by three complementary roles working together toward publication-ready, clinically significant research.

ML Fellow (rising juniors/seniors)

Technical Lead

- Lead technical design & coding
- Ensure reproducibility
- Own day-to-day execution
- Present at conferences

Med Mentor (You!)

Clinical Guide

- Provide clinical framing
- Guide project scoping
- Ensure clinical rigor
- Support manuscript prep

Faculty Mentor (Your PI)

Principal Investigator

- Provide clinical question
- Approve methodology
- Ensure IRB compliance
- Support publication

Two Pathways: Learn AI Skills or Leverage Existing Skills

Pathway 1

No AI Skills? Learn & Lead Clinically

Learn technical & AI skills through mentorship and hands-on work

Provide clinical direction and clinical input to the project

Develop leadership skills: Clinical judgment + technical literacy

Perfect for: Medical students wanting to learn AI applications

Pathway 2

Have AI Skills? Contribute Programming

Do programming & technical work on the AI project

Apply clinical context to your technical contributions

Develop leadership skills: Mentoring fellows and guiding strategy

Perfect for: Medical students with technical background

Research Incubator: Accelerating Clinical AI Publications

The triad model and mentorship create a multiplier effect that accelerates publication and creates authorship opportunities for all participants.

How It Works

→ **Fellows:** programs models and provides technical guidance

→ **Med Mentor:** promote clinical guidance and manuscript support

→ **Faculty Mentor:** ensure rigor and publication strategy

→ **Result:** Faster publication timeline with mentorship at every step

Authorship Opportunities

→ **Fellow:** First or lead author on publications

→ **Med Mentor:** Co-first-author on their fellow's clinical papers and manuscripts and potentially co-authorship on other projects

→ **Faculty Mentor:** Senior author and publication advocate

→ **Chapter Bylaws:** ensure fair attribution and recognition

Publication Impact

→ **Accelerated timeline:** From project to publication **in one summer** is **very realistic** in the AI research space

→ **High-impact venues:** Peer-reviewed journals and conferences

→ **Career advantage:** Publications strengthen residency applications

Project Timeline

March

Project Incubation

Mentors refine their hypotheses with senior faculty. We ensure every project is viable, high-impact, and ready for a technical partner.

March – April

Fellowship Selection

We open the doors to engineers. Medical students interview and select their own technical counterparts, building a team perfectly matched for the problem at hand.

April – June

The Onramp

Teams secure IRBs, data access, and credentials now so that when summer starts, you can focus 100% on building.

June – August

10-Week Research Sprint

Build the model, run experiments, analyze results, and write the paper. This is the core of the fellowship.

Fellowship curriculum: 1 hr/week

- **30 minutes of small group journal club discussion about topics in medical AI foundations, model safety, ethics, commercialization, etc**
- **30 minutes of project updates**

September

Present your work

Share your findings at the Tensor Lab Symposium and submit your paper for publication

FAQ

- ***Do medical students need to know how to code?***
  - **No.** We match you with a Machine Learning fellow who can handle all the coding.
- ***How much time do medical students spend per week?***
  - **About 2-4 hours per week.** You'll meet with your fellow regularly to review progress and guide the direction of the research, and participate in weekly meetings with the whole cohort.
- ***Do ML fellows need a medical background?***
  - **No.** ML fellows will come with strong Python/ML skills to support you and your project. You'll teach them the relevant medical background.
- ***How much time do ML fellows spend per week?***
  - They'll spend **15-20 hours per week**.
- ***Is this remote or in-person?***
  - **Mostly remote.** You'll meet with your team over Zoom/Discord, although our chapter may organize optional in-person sessions.
- ***Is this paid?***
  - **No.** Neither the ML fellow nor the med student will receive additional funding at this time.
- ***Who owns the code and research?***
  - The code will be open-sourced, and all team members share authorship. Patient data stays under the PI's institutional controls.
- *What do I need to apply?*
  - **You need a PI who can add an unpaid external volunteer to your IRB-compatible project.**
  - **You do NOT need an AI project idea.** Tensor Lab directors can help you brainstorm potential ideas.

Pilot Program 2025: a success story

Pilot Program 2025: Strong Mentorship Network

Network Strength

Mentored by **established leaders in clinical AI and data science** from top-tier institutions.

Featured Institutions

- University of California, San Francisco (UCSF)
- University of Maryland, School of Medicine

> **Unparalleled networking with PIs and department heads shaping the future of medicine.**

Featured Mentors

Dr. Andreas Rauschecker, MD, PhD

Co-Director, Center of Intelligent Imaging, UCSF. Leader in advanced medical imaging and AI-driven diagnostics.

Dr. Stefano Bini, MD

Maria Manetti Shrem Endowed Professor, Orthopedic Surgery, UCSF. Expert in surgical innovation and outcomes research.

Dr. Jonas De Almeida, PhD

Director of data science and senior investigator, National Cancer Institute. Leader in clinical AI research and data science.

Dr. Bradley A. Maron, MD

Director of Pulmonary Hypertension Program & Co-Director of UM-IHC, University of Maryland. Director of Scientific Operations.

Plus faculty mentors across oncology, radiology, and other specialties.

Pilot Program 2025: Demonstrated Success

Our inaugural pilot program demonstrated the power of the Tensor Lab model with exceptional outcomes across all metrics.

Talent Pool

200+

- ✓ Applications received
- ✓ $50K NIH grant winners
- ✓ NASA, Google, Apple internships

Research Output

8/10

- ✓ Fellows presented posters
- ✓ 10 total in cohort
- ✓ 80% presentation rate

Publication Impact

4+

- ✓ In submission to journals/conferences
- ✓ High-impact venues
- ✓ Rigorous, reproducible research

Benchmarking / ICU Data NCI

LLMs vs. Traditional ML

FELLOW: JIALONG • MENTOR: AARON GE

Compared the performance of LLM-based semantic models (using LLM embeddings) against traditional ML models (XGBoost) across different clinical prediction tasks.

Result: Text representation critically affects performance; traditional ML excels at physiological data.

Large Language Models vs. Traditional Machine Learning for Clinical Risk Prediction: A Mechanistic Analysis of When Statistical and Semantic Models Fail

Aaron Ge[1], Jialong Wu[3], Jonas De Almeida[2], Bradley Maron[1]*

[1] Institute of Health Computing, University of Maryland, School of Medicine, Baltimore, MD, USA

[2] Division of Cancer Epidemiology and Genetics, National Cancer Institute, National Institutes of Health, Maryland, USA

[3] Department of Computer Science, Whiting School of Engineering, Johns Hopkins University, Baltimore, MD, USA

BACKGROUND

- Clinical prediction can be built on two paradigms: statistical models using structured numeric features versus semantic models that encode patient text with large language model (LLM) embeddings.
- It remains unclear when each paradigm excels or fails, and how input design (representation) and prompting influence performance.
- Understanding the mechanisms baehind disagreements between paradigms is critical for safe deployment in healthcare.

OBJECTIVE

Predict 6 different binary endpoints (in-hospital mortality,) using 458 engineered clinical features derived from the first 24 hours of ICU stay.

- **H1 LLM embeddings vs. XGBoost**: Optimized LLM embedding model will perform similarly to an optimized XGBoost model.
- **H2 Hybrid Model Benefit**: Models make different mistakes, so combining these models will improve performance.
- **H3 LLM Sensitivity**: Semantic system performance is sensitive to text serialization and LLM prompt task-alignment.

METHODS

1. **Data Preparation**
   - **Dataset:** ICU patient data with 458 clinical features from first 24 hours of admission.
   - **Prediction Tasks:** In-hospital mortality, Vasopressor use, Mechanical ventilation, ICU length of stay >3 days, Total length of stay >7 days, 30-day readmission.
2. **Model Development**
   - Numerical Models (NM)
   - Semantic Models (SM)
     - **Text Serialization:** F1: Raw feature values, F2: Key-value pairs, F3: Narrative clinical summary.
     - **LLM Embeddings:** Multiple models compared: General-purpose LLMs, Medical-specific LLMs.
3. **Hybrid Model Fusion:**
   - **25+ fusion strategies systematically tested**
4. **Evaluation Framework**
   - **Primary Metrics:** AUROC, AUPRC.
   - **Statistical Analysis:** Stratified bootstrap (1000 iterations), 95% confidence intervals.

RESULTS

H1 LLM embeddings vs. XGBoost: Human-engineered features still beat semantic embeddings.There is one notable exception.The text-embedding-005 model using the F3 representation did show a slight performance advantage over the XGBoost baseline for the 30-Day Readmission task.

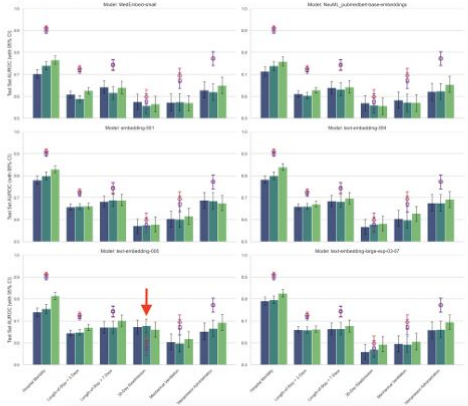

H2 Hybrid Model Benefit: Synergy is observed in 4 out of 6 tasks. For the task where the base models performed best (Mortality), fusion failed entirely.

Prediction Task	Best Baseline Model	Best Baseline AUROC (95% CI)	Best Fusion Strategy	Best Fusion AUROC (95% CI)	Absolute Change
In-Hospital Mortality	Numerical	0.9107 (0.8998-0.9213)	No Synergy Observed	0.9100 (0.8968-0.9204)	-0.0007
30-Day Readmission	Semantic	0.6776 (0.6539-0.7101)	No Synergy Observed	0.6712 (0.6412-0.7020)	-0.0064
Vasopressor Intervention	**Numerical**	**0.7724 (0.7495-0.8025)**	**Neural Network**	**0.7862 (0.7548-0.8176)**	**+0.0138***
Ventilation Intervention	Numerical	0.6939 (0.6602-0.7271)	Hierarchical	0.7023 (0.6690-0.7356)	+0.0084
Length of Stay > 3 Days	Numerical	0.7265 (0.7140-0.7397)	Dynamic Weight	0.7324 (0.7198-0.7452)	+0.0059
Length of Stay > 7 Days	Numerical	0.7438 (0.7203-0.7682)	Ensemble	0.7499 (0.7268-0.7747)	+0.0061

H3 Prompt Sensitivity: Prompting is a Surprisingly Ineffective Strategy. There is no evidence that complex, task-aligned prompts provide any statistically significant benefit over a simple Control prompt (P0).

Figure 3: AUROC Improvement Over Best Baseline for Hospital Mortality

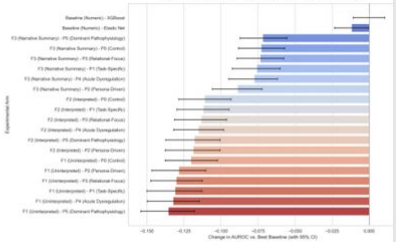

CONCLUSIONS

- **Numerical Models Remain Strong:** Traditional numerical models outperform semantic embedding models for most prediction tasks, despite advanced prompt engineering and diverse data serialization techniques.
- **Narrative Data Shows Promise:** The format of data for LLMs is crucial. Succinct representations (F3) are more effective in capturing clinically relevant information than structured value lists.
- **SM and NM Have Distinct Failure Modes:** Synergy is possible, but it is conditional and not a general property of combining these two modalities. Numerical models do better on patients with invasive hemodynamic monitoring. Semantic models do better on patients with AKI.

Head & Neck Cancer UMSOM

RFS Prediction via PET/CT

FELLOW: KOUSHIK • MENTOR: LAURA CHEN

Investigated whether a simplified multimodal model using only PET/CT and clinical data could achieve comparable RFS prediction to expert-annotated models.

Result: Simpler PET/CT model achieved comparable performance and is more generalizable.

The Tensor Lab for Computational Medicine

Could Simplified Deep Learning be More Generalizable? Predicting Recurrence-Free Survival in Head and Neck Cancer With and Without Tumor Segmentation

Johns Hopkins Whiting School of Engineering

Koushik Rameshbabu, Laura Chen, Aaron Ge, Dr. Grace Snow

UCSF

University of Maryland School of Medicine

Introduction

- Head and Neck Cancer (HNC) accounts for over 900,000 new cases and 400,000 deaths annually worldwide, making it one of the most prevalent malignancies.
- Despite advances in radiotherapy and chemotherapy, recurrence rates remain high (30–50%), emphasizing the need for accurate recurrence-free survival (RFS) prediction.
- RFS prediction can help tailor care where low-risk patients may benefit from reduced chemotherapy or surveillance while high-risk patients can receive closer monitoring and timely intervention.
- Many deep learning models rely on segmentation-derived tumor features which require expert annotation. This study investigates whether a simplified multimodal model using only PET/CT and clinical data can achieve comparable predictive performance while improving generalizability.

Dataset Description

- The dataset was obtained from the HECKTOR 2025 Challenge, including 651 FDG-PET/CT scans and clinical data from 7 international centers.
- PET and CT images were co-registered, resampled to 1 mm, cropped around tumors, and z-score normalized.
- Clinical variables included age, gender, tobacco, alcohol, performance status, M-stage, HPV status, treatment, and all were standardized.
- Performance status was grouped as Functional (ECOG 0–1 / Karnofsky 80–100) or Restricted (ECOG 2–4).
- Tumor features such as volume, surface area, compactness, and sphericity were extracted from segmentation masks and min–max normalized.
- The cohort had a 20.3% event rate, with 132 relapses and 519 censored cases, providing a balanced 4:1 ratio suitable for survival analysis.

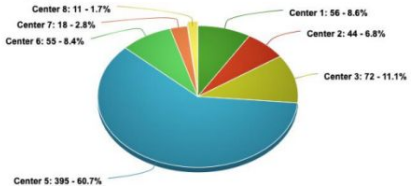

Methodology

We developed a four-stage multimodal deep learning pipeline integrating CT, PET, and clinical data for recurrence-free survival prediction, combining feature extraction, fusion, and ensemble survival modeling to capture both anatomical and clinical prognostic information.

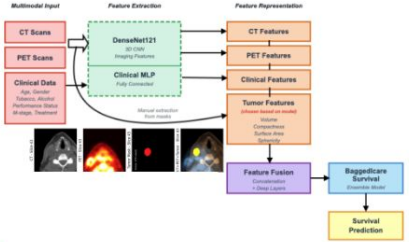

- **Stage 1** – *Multimodal Input*: Integrated co-registered FDG-PET/CT scans with key clinical variables (age, gender, tobacco, alcohol, performance status, M-stage, treatment, HPV status).
- **Stage 2** – *Feature Extraction*: Imaging features were derived using DenseNet121 (3D CNN) and clinical features via a fully connected MLP, with tumor features manually extracted from segmentation masks.
- **Stage 3** – *Feature Fusion & Representation*: Extracted CT, PET, clinical, and tumor features were concatenated to form a unified deep representation.
- **Stage 4** – *Survival Prediction*: The fused features were used in a Bagged ICARE survival ensemble for patient-level recurrence-free survival estimation.
- A stratified center-wise split ensured ~20% of patients per center in the test set (130 total) and 5-fold cross-validation on 521 remaining patients for robust evaluation.

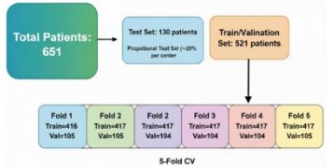

Results

- Model performance was evaluated using the Concordance Index (C-index) and Kaplan–Meier (KM) survival curves.
- The C-index measures how well the model ranks patients by recurrence risk where higher values indicate better agreement between predicted and actual survival orderings.
- The Kaplan–Meier curves visualize patient survival over time, grouped into low, medium, and high predicted-risk tertiles. Greater vertical separation between curves reflects stronger discrimination.

Model type	Mean Validation C-index	Mean Test C-index	Best Validation C-index	Best Test C-index
With Tumor Features	0.6644	0.6603	0.7248	0.6791
Without Tumor Features	0.6371	0.5533	0.6968	0.6367

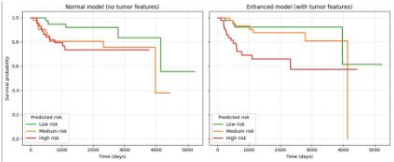

- ✓ The enhanced model (with tumor features) achieved higher mean C-index values and clearer high-risk separation on KM plots, showing improved early recurrence prediction.
- ✓ The simpler model (without tumor features) remained competitive and more generalizable across centers despite slightly lower C-index performance.

Conclusion

Tumor features enhanced recurrence-free survival prediction and improved the C-index, confirming their prognostic value. However, the simpler PET/CT + clinical model achieved comparable performance and remains more generalizable for real-world clinical deployment where segmentation data may be unavailable.

Cardiology / Readmission UCSF

Predicting Readmissions

FELLOW: SHRAVYA • MENTOR: GRIFFITH HUGHES

Developed an attention-based deep learning model (BioClinicalBERT) to predict 30-day readmissions from free-text discharge summaries.

Result: Achieved AUROC 0.730 with interpretable predictions via SHAP.

Predicting 30-Day Hospital Readmissions Using Attention-Based Deep Learning and Explainable AI on Clinical Discharge Summaries

The Tensor Lab for Computational Medicine

Shravya Dasu (Student Fellow), Griffith Hughes (Medical Student Mentor), Dr. Geoff Tison (Faculty Mentor)
Affiliations: The Tensor Lab for Computational Medicine

https://github.com/ShravyaD-coder/30-Day-Readmission-XAI.git

Introduction

Hospital readmissions within 30 days affect one in five Medicare beneficiaries and cost over $26 billion annually. Current prediction models rely on structured data and overlook rich clinical narratives in discharge summaries that contain modifiable risk factors. While MIMIC-III has been widely used for readmission prediction research, MIMIC-IV clinical notes remain underexplored.

Objectives

To develop and validate an attention-based deep learning model that predicts 30-day hospital readmissions directly from discharge summaries, with explainable AI techniques to identify clinically relevant risk factors

Materials

- **Data**

MIMIC-IV Note: Deidentified free-text clinical notes from patients at the Beth Israel Deaconess Medical Center in Boston. Contains **331,794** de-identified **discharge summaries** from **145,915 patients**

MIMIC-IV Hospital: Contains **546,028** unique **hospitalizations** for **223,452** unique **individuals** admitted to the BIDMC emergency department.

- **Pre-processing**

Data Integration: Merged clinical notes with hospital admission records by patient ID

Outcome Definition: Calculated 30-day readmission from discharge-to-admission dates; excluded deaths and missing dates

Text Cleaning: Removed extra whitespace, special characters, and de-identified artifacts

Data Split: Patient-level grouping with stratified split (**Train: 72%, Validation: 8%, Test: 20%**) to prevent data leakage and maintain readmission proportion across sets

Training Comparison: 10% subset (3 epochs) vs. full dataset (1 epoch

Figure 1. Readmission Category Proportions in Data Splits

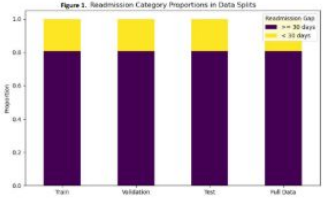

- **Model Architecture:**

Base Model: BioClinicalBERT, a BERT model trained on PubMed articles and the entirety of MIMIC-III notes

- **Evaluation Metrics:**

AUPRC, AUROC, Accuracy

Methodology

Results

```
Readmission Probability: 77.3%
Prediction: Readmitted

[Chunk 4] Attention Weight: 0.428
--------------------------------------------------------------------------------
---
##ine hcl 100 mg tablet sig: one (1) tablet po daily (daily). 2. folic acid 1
mg tablet sig: one (1) tablet po daily (daily). 3. amitriptyline
... (truncated)

[Chunk 3] Attention Weight: 0.273
--------------------------------------------------------------------------------
---
##bral compression fracture. brief hospital course: abdominal distention /
pain: she was treated empirically due to concern for spontaneous bacterial
peritonitis with ceftriaxone 2g x 1. a diagnostic paracentesis was performed
... (truncated)

[Chunk 1] Attention Weight: 0.242
--------------------------------------------------------------------------------
---
history of present illness : with recently diagnosed alcoholic hepatitis,
persistent ascites, and persistent fevers and leukocytosis which have been
atributed to her hepatitis who presented to today with worsening abdominal
distention, pain, and persistent fever. she denies
... (truncated)
```

Figure 2: Chunks of note having the highest attention weight contributing to high probability of readmission

Table 1: Performance Comparison of 2 fine-tuning approaches

Metric	10% Data (3 Epochs)	Full Data (1 Epoch)
AUROC	0.692	0.730
AUPRC	0.372	0.415
Accuracy	64%	73%

```
Top chunks INCREASING readmission risk (positive SHAP):

1. SHAP: +0.288
   ##ine hcl 100 mg tablet sig : one ( 1 ) tablet po daily ( daily ). 2. folic acid 1 mg tablet sig : one ( 1 ) tablet

2. SHAP: +0.274
   ##bral compression fracture. brief hospital course : abdominal distention / pain : she was treated empirically due

3. SHAP: +0.187
   name :. service : medicine patient recorded as having no known allergies to drugs chief complaint : abdominal diste
```

Figure 3: SHAP on the high attention weight chunks to see how they contribute to readmission prediction

Conclusion, Limitations and Future Work

- **Conclusion:** The attention-based model achieves AUROC 0.730 with interpretable predictions, identifying specific discharge summary sections containing modifiable risk factors
- **Key-Takeaways:** Training on full data for 1 epoch outperformed 10% data for 3 epochs, demonstrating that data quantity is more impactful than additional training epochs for this task
- **Limitations:** Since the data was collected from a single center, it limits the scope of generalizability of our data and the high class imbalance impacts our precision negatively
- **Future Work:** With the availability of sufficient computing resources, the entire dataset will be used for analysis and if possible, clinical notes from different medical centers will be used for evaluation to test generizabilty of the approach

Acknowledgement and Contact

I gratefully acknowledge the MIMIC-IV database provided by the MIT Laboratory for Computational Physiology, Beth Israel Deaconess Medical Center, and PhysioNet, which made this research possible. I thank Griffith Hughes for his invaluable guidance and support throughout this project. I am deeply grateful to The Tensor Lab for Computational Medicine for providing this research opportunity and fellowship support.

Fellow Stories

“

Tensor Lab showed me meaningful challenges that combine both clinical and technical domains. As I prepare for my PhD applications, this experience was invaluable in clearly defining my research direction.

Jialong
Research Fellow

“

Tensor Lab taught me to conceptualize and develop applications for LLMs that require significant domain knowledge.

Ashwin
Research Fellow

“

The guidance from the faculty and mentors was exceptional... I wouldn’t have reached our final publication goals without their help.

Koushik
Research Fellow

“

Collaborating with a global cohort... provided me with a network of contacts for future projects.

Juan Arturo
Research Fellow

Sounds great!
How do I apply?

Use this form!

Apply by Sunday 4/5!

https://forms.gle/gNbG7Gegxyb5oALL8

THE TENSOR LAB
FOR COMPUTATIONAL MEDICINE

The Tensor Lab — Machine Learning Fellow Request (PRISM/Other Summer Research Fellowship applicants)

Request an ML Fellow

Use this form to request a machine learning fellow for your research project. ML Fellows are advanced undergraduate and graduate students in computer science from UMCP, Johns Hopkins, and other universities who are matched with your team to integrate ML methods into an existing project or develop a new research component from scratch.

About the Program

Tensor Lab is a research fellowship that pairs ML engineers with medical students and senior physicians to tackle high impact clinical problems. Over a 10 week summer sprint, each multidisciplinary team works with real world medical datasets to build and refine models. The program culminates in academic dissemination of the team's work, with the goal of advancing clinical decision making and patient outcomes.

Your Role

As the requesting medical student, you will serve as the lead of your triad, guiding the project's clinical direction while your ML Fellow brings the technical expertise.

Use this!

THE TENSOR LAB
FOR COMPUTATIONAL MEDICINE

Medical Student Interest Form

Thank you for your interest in joining The Tensor Lab for Computational Medicine. We are looking for medical students passionate about advancing medicine with machine learning.

The Tensor Lab fellowship is a research program for machine learning engineers, putting them in teams with medical students and senior physicians to solve high-impact clinical challenges. Over a 10-week summer sprint, these multidisciplinary triads collaborate to build and improve models using real-world medical datasets. The program culminates in the academic dissemination of the team's achievements to enhance clinical decision-making and patient care.

When you join, you will be the leader of this triad for your research project.

Applications for the Medical Mentor position are due **March 6th at 11:59 PM PST**.

Please reach out to us or see our website for more details: www.thetensorlab.org

Not the form on our website!

Thank you for coming!

Apply by Sunday 4/5!

Attendance!